

An Automated Framework for Dataset Quality Assessment and Data Leakage Detection in Machine Learning

Jaime Cano Morano
Illinois Institute of Technology

June 2026

Abstract. Data quality issues and data leakage are among the most common and damaging problems in applied machine learning, yet most practitioners lack systematic tools to detect them before model training. We present **ml-framework**, an open-source Python framework that automates dataset assessment across six dimensions: data quality, leakage detection, feature analysis, statistical sufficiency, distributional drift, and impact quantification. The framework introduces a novel *unified leakage risk score* that combines Pearson correlation, normalised mutual information, and per-feature performance inflation into a single interpretable metric, and a *semantic leakage module* that employs a large language model (GPT-4o-mini) to identify implicit future-information features that purely statistical methods cannot detect. Experiments on five real-world and six synthetic datasets demonstrate that ml-framework detects 4/4 constructed leakage scenarios, covers 29 distinct checks, and outperforms ydata-profiling, Deepchecks, and Great Expectations in both feature coverage and leakage detection rate. All code, datasets, and benchmarks are available at <https://github.com/jaimecano12/ml-framework>.

1 Introduction

Modern machine learning pipelines depend critically on the quality of their training data, yet practitioners routinely overlook systematic pre-training assessment. Kaufman et al. [1] identified data leakage as “one of the most important problems in data mining”, describing scenarios where models appear to achieve near-perfect accuracy on validation sets but fail catastrophically in production. Sculley et al. [2] further documented how undiscovered data quality problems accumulate as *hidden technical debt* that is expensive to repay once models are deployed.

Despite this awareness, the existing ecosystem of data validation tools focuses primarily on schema validation and exploratory analysis rather than ML-specific issues. Libraries such as ydata-profiling [4], Great Expectations [5], and Deepchecks [6] provide valuable engineering utilities but offer little support for detecting temporal leakage, quantifying performance inflation caused by leaked features, or producing actionable readiness scores tailored to a specific prediction task.

To illustrate the scope of the problem, consider three scenarios encountered during our evaluation. First, the Titanic dataset from OpenML contains a `boat` column (lifeboat assignment number) with Pearson $|r| = 0.97$ against the survival target—a feature that encodes the outcome and would never be available at prediction time. A naive model trained on this dataset achieves near-perfect accuracy not because it has learned meaningful survival patterns, but because it has memorised a post-hoc label. Second, a synthetic credit-scoring dataset with a noisy proxy feature ($r \approx 0.97$) achieves baseline cross-validated accuracy of 1.000; removing the leaky feature drops accuracy to 0.952, a $\Delta = -0.048$ drop that reveals entirely artificially inflated performance. Third, medical datasets frequently contain features such as `discharge_code`

or `mortality_score` whose leakage is invisible to correlation and mutual information analyses but obvious from the feature name semantics—only a language-aware analysis can flag them.

These examples motivate a framework that goes beyond simple schema validation and combines statistical, information-theoretic, and language-based detection methods into a single, actionable pipeline.

This paper presents **ml-framework**, a modular Python framework designed to address these gaps. Our main contributions are:

1. A unified pipeline of **21 automated checks** spanning six analysis dimensions, accessible via CLI, Python SDK, and a Streamlit dashboard.
2. A **unified leakage risk score** that aggregates correlation, mutual information, and per-feature performance inflation into a single $[0, 1]$ metric, providing a more principled signal than any individual heuristic.
3. A **semantic leakage analysis module** driven by GPT-4o-mini that detects implicit future-information features based on feature names and dataset descriptions—leakage that statistical methods cannot detect.
4. A **quantitative benchmark** on four installed tools, showing that ml-framework covers $3.2\times$ more checks and detects $4\times$ more leakage scenarios than the best-performing alternative.

2 Related Work

Data leakage. Kaufman et al. [1] formalised data leakage as any information about the label that is incorporated into the model but would not be available at prediction time. They categorise leakage into target leakage (features directly derived from the target), train/test contamination (validation data influencing training), and temporal leakage (future information used as a feature). Our framework operationalises all three categories with dedicated checks and adds a fourth: *ID-column leakage*, where high-cardinality identifiers enable memorisation.

Data quality tools. **ydata-profiling** [4] generates exploratory analysis reports with statistics, correlation matrices, and distribution plots but provides no ML-specific leakage detection or readiness scoring. **Great Expectations** [5] is a data contract framework that validates user-defined schema expectations; it requires manual rule authoring and has no concept of target leakage. **Deepchecks** [6] is the closest competitor, offering train/test integrity suites and some drift detection, but it covers only 11 of the 29 checks we identify as relevant (Section 6).

Data-centric AI. The data-centric AI movement [16], championed by Ng, argues that improving data quality is more impactful than tuning model architectures for most production tasks. Breck et al. [3] formalised this at scale with TFX Data Validation (TFDV), which monitors feature statistics and schema conformance in production ML pipelines. Our framework is complementary to TFDV: while TFDV excels at production monitoring and schema enforcement, ml-framework is designed for the *pre-training* assessment phase, where leakage detection and readiness scoring matter most.

LLMs for data quality. The use of large language models for structured data tasks has grown rapidly. Narayan et al. [17] showed that GPT-3 can perform entity matching and data imputation with competitive accuracy. Sui et al. [18] explored LLM-assisted table understanding. Our semantic leakage module extends this line of work by tasking GPT-4o-mini with a novel role: identifying features whose *names* implicitly encode future or post-hoc information—a task that requires domain knowledge and semantic reasoning unavailable to purely statistical methods.

Mutual information for feature selection. Kraskov et al. [9] introduced non-parametric MI estimation used in our unified risk score. Ross [10] extended this to mixed continuous-discrete variables, matching our use case of heterogeneous tabular data.

Drift detection. Rabanser et al. [11] evaluated two-sample statistical tests for covariate shift detection. Our covariate drift module employs the Kolmogorov-Smirnov (KS) test with Bonferroni correction alongside the Population Stability Index (PSI) [12], a complementary binning-based measure widely used in credit risk modelling.

3 System Architecture

3.1 Design Principles

ml-framework is designed around three principles: *modularity* (each check is an independent function), *composability* (checks are orchestrated by a single pipeline with a shared configuration schema), and *extensibility* (a plugin system allows third-party checks to be registered at runtime).

3.2 Pipeline Overview

Figure 1 depicts the full analysis pipeline. A dataset and a YAML configuration file are the only required inputs. The pipeline executes six analysis phases in parallel; results are aggregated into a `FrameworkReport` and passed to the scoring and recommendation engines.




Figure 1. ml-framework analysis pipeline. A dataset and a configuration file flow through six parallel analysis phases, are aggregated into a `FrameworkReport`, scored (0–100, grade A–F), and exported as HTML, JSON, or via Python SDK.

3.3 User Interfaces

ml-framework exposes three interfaces targeting different user profiles. The **command-line interface (CLI)** accepts a dataset path and configuration file, executes the full pipeline, and writes an HTML report to a configurable output directory—suitable for scripting and CI/CD integration. The **Python SDK (DatasetChecker)** provides programmatic access for Jupyter notebooks and automated model-training pipelines. Listing 1 shows a complete end-to-end usage example. The **Streamlit web application** offers a no-code interface for exploratory use: users upload a CSV, Parquet, or Excel file, configure checks via sidebar widgets, and navigate results across seven tabs (Summary, Quality, Leakage, Features, Sufficiency, Drift, and Recommendations), with one-click download of the HTML report and JSON export.

```
1 from src.checker import DatasetChecker
```

```

2
3 checker = DatasetChecker("configs/config.yaml")
4 report  = checker.run("data/titanic.csv",
5                       target_col="survived")
6
7 # Readiness score and grade
8 print(f"{checker.score}/100   grade={checker.grade}")
9
10 # Failed checks with severity
11 for r in report.failed_checks():
12     print(f"[{r.severity}] {r.check_name}: {r.message}")
13
14 # Prioritised recommendations
15 for rec in report.recommendations[:3]:
16     print(f"[{rec.priority}] {rec.action}")
17
18 # Export
19 checker.save_report("reports/")          # HTML + JSON
20 d = checker.to_dict()                    # JSON-serialisable dict

```

Listing 1. Python SDK: complete usage example.

3.4 Module Structure

Table 1 summarises the 14 source modules. The shared dataclass hierarchy (CheckResult, FrameworkReport, ReadinessScore) defined in `utils.py` is used by all modules.

Table 1. Source modules and primary responsibilities.

Module	Responsibility
quality_checks	Missing values, duplicates, outliers, class imbalance, constant/low-variance
leakage_checks	Target leakage, train/test overlap, temporal, ID column, unified risk score
feature_analysis	Feature correlation, MI relevance, distribution shape
sufficiency	Sample size, class support, CV stability, feature-to-sample ratio
drift_checks	KS+PSI covariate drift; χ^2 label drift
impact_analysis	Baseline vs. cleaned cross-validation delta
scoring	Weighted readiness score (0–100) and letter grade (A–F)
recommendations	20 handlers mapping checks to code-level suggestions
semantic_leakage	GPT-4o-mini semantic leakage analysis
reporting	Jinja2 HTML report with embedded matplotlib figures
checker	DatasetChecker Python SDK facade
plugins	@register_check decorator and module loader
config	YAML loader with deep-merge and schema validation

4 Core Methodology

4.1 Quality Checks

Six checks form the data quality dimension (weight $w_q = 0.25$):

Missing values. Columns with NaN rate exceeding τ_{miss} (default 5%) are flagged as warning (5–20%) or error (>20%).

Duplicates. Exact row duplicates indicate data collection errors or potential train/test contamination.

Outliers. Both IQR and z -score methods are supported. IQR flags values outside $[Q_1 - k \cdot \text{IQR}, Q_3 + k \cdot \text{IQR}]$ with $k = 3.0$ by default.

Class imbalance. Minority class ratio $r = n_{\text{min}}/N < 0.10$ triggers a warning; $r < 0.05$ triggers an error.

Constant and low-variance features. Low-variance detection uses the coefficient of variation $CV = \sigma/|\mu|$ rather than min-max normalised variance, which always maps to $[0, 1]$ and fails to detect near-constant features.

4.2 Leakage Detection

4.2.1 Classical Checks

Target leakage. For numeric features we compute Pearson $|r|$; for categorical features we use Cramér’s V . Features with association ≥ 0.95 are flagged as errors.

Train/test overlap. A stratified split is simulated and duplicate rows spanning the two halves are counted. An overlap rate $> 5\%$ is an error.

Temporal leakage. If a date column is configured, monotonic chronological order is verified. An unsorted temporal sequence means a naive sequential split would mix past and future observations.

ID-column leakage. String and integer columns with unique-value ratio $\geq 95\%$ are flagged. Float columns are excluded, as continuous variables naturally exhibit high cardinality.

4.2.2 Unified Leakage Risk Score

The four classical checks are heuristic thresholds applied independently. We introduce a *unified leakage risk score* $\mathcal{L}(f)$ that aggregates three complementary signals for each feature f :

$$\mathcal{L}(f) = w_c \cdot \rho(f) + w_m \cdot \tilde{I}(f; y) + w_p \cdot \pi(f) \quad (1)$$

where:

- $\rho(f) \in [0, 1]$: Pearson $|r|$ (numeric) or Cramér’s V (categorical) between feature f and target y .
- $\tilde{I}(f; y) = I(f; y) / \max_{f'} I(f'; y) \in [0, 1]$: mutual information normalised by the cross-feature maximum, capturing non-linear associations that correlation misses.
- $\pi(f) \in [0, 1]$: *performance inflation*, defined as

$$\pi(f) = \frac{\max(0, A_f - A_{\text{base}})}{1 - A_{\text{base}}}$$

where A_f is the 3-fold cross-validated accuracy of a logistic regression trained on f alone, and $A_{\text{base}} = \max_c p(y = c)$ is the majority-class baseline.

Default weights are $w_c = w_m = 0.35$, $w_p = 0.30$; all are configurable via YAML. Features with $\mathcal{L}(f) \geq 0.7$ are flagged as warnings; ≥ 0.9 as errors.

Algorithm 1 presents the full computation procedure. A key design decision is the use of a single-feature logistic regression (rather than a more complex model) for $\pi(f)$: it is fast, interpretable, and provides a linear upper bound on the predictive signal that a feature can contribute independently of all others.

4.3 Feature Analysis

Feature correlation. Pairs with Pearson $|r| \geq 0.90$ indicate redundant information inflating effective dimensionality.

Feature relevance. Normalised mutual information < 0.01 flags features contributing negligible information about the target.

Distribution shape. $|\text{skewness}| > 2$ or excess kurtosis > 7 signal non-Gaussian distributions that may violate model assumptions.

4.4 Statistical Sufficiency

Four checks assess whether the dataset is statistically adequate for the modelling task: minimum sample size ($n \geq 100$ and $n/p \geq 50$), minimum class support (30 samples per class), cross-validation stability (CV std ≤ 0.10), and feature-to-sample ratio ($p/n \leq 0.10$).

Algorithm 1 Unified Leakage Risk Score computation.

Require: DataFrame D , target column y , weights w_c, w_m, w_p

- 1: Encode categoricals; impute missing values with column means
- 2: $A_{\text{base}} \leftarrow \max_c p(y=c)$ ▷ majority-class baseline
- 3: **for** each feature $f \in D \setminus \{y\}$ **do**
- 4: $\rho(f) \leftarrow$ Pearson $|r|$ or Cramér’s V between f and y
- 5: $I(f; y) \leftarrow$ MutualInfoClassif(f, y)
- 6: $A_f \leftarrow$ mean 3-fold CV accuracy of LR trained on $\{f\}$ alone
- 7: $\pi(f) \leftarrow \max(0, A_f - A_{\text{base}}) / (1 - A_{\text{base}})$
- 8: $\mathcal{L}(f) \leftarrow w_c \cdot \rho(f) + w_m \cdot \tilde{I}(f; y) + w_p \cdot \pi(f)$
- 9: **end for**
- 10: Normalise $\tilde{I}(f; y) \leftarrow I(f; y) / \max_{f'} I(f'; y)$
- 11: **return** $\{f : \mathcal{L}(f) \geq \tau\}$ flagged as high-risk leakage

4.5 Drift Detection

The dataset is split chronologically into reference and current halves. For each numeric feature we apply the two-sample KS test and compute PSI:

$$\text{PSI} = \sum_{b=1}^B (p_b^{\text{curr}} - p_b^{\text{ref}}) \ln \frac{p_b^{\text{curr}}}{p_b^{\text{ref}}} \quad (2)$$

with adaptive binning (minimum 20 observations per bin) and Bonferroni correction across all features. Label drift is assessed with a χ^2 test on the target distribution.

4.6 Impact Analysis

The framework trains classifiers on two dataset variants: the *baseline* (all features) and the *cleaned* (problematic columns removed, duplicates dropped). The performance delta $\Delta = A_{\text{clean}} - A_{\text{base}}$ quantifies how much a model benefits from leaky or low-quality features. A large negative Δ confirms that the baseline accuracy was inflated by data issues rather than genuine predictive signal.

4.7 Readiness Score

The composite readiness score integrates the four main analysis dimensions:

$$S = 100 - (w_q \cdot P_q + w_l \cdot P_l + w_f \cdot P_f + w_s \cdot P_s) \quad (3)$$

where P_d is the total penalty for dimension d (errors cost 15 points per check; warnings cost 5 points), and the weights are $w_q=0.25, w_l=0.35, w_f=0.25, w_s=0.15$. Leakage penalties are further amplified by a $1.5\times$ multiplier. Letter grades: A (≥ 85), B (≥ 70), C (≥ 55), D (≥ 40), F (< 40).

Weight rationale. Leakage receives the highest weight ($w_l=0.35, 1.5\times$ multiplier) because it is a silent failure mode: a model trained on leaky data appears to perform well during validation but generalises poorly in production [1]. Quality and feature issues ($w_q=w_f=0.25$) are equally weighted as they are typically correctable. Sufficiency ($w_s=0.15$) contributes least because sample-size adequacy depends strongly on the chosen learner. Penalties were calibrated so a single confirmed leaky feature drops a dataset from grade A to grade B. All parameters are overridable via the `scoring` block in `config.yaml`.

4.8 LLM Semantic Leakage Analysis

Statistical methods cannot detect *semantic* leakage, where a feature name itself reveals post-hoc or future information (e.g. `discharge_code`, `future_usage`). We address this with an optional module that submits feature names, sample values, and a user-provided dataset description to GPT-4o-mini via Azure

OpenAI. The model classifies each feature into one of five risk levels (none, low, medium, high) and four leakage types (temporal, proxy, post_hoc, indirect). The module degrades gracefully when no API credentials are configured.

Quantitative evaluation. To evaluate the module systematically, we constructed a benchmark of 30 manually labelled features drawn from five realistic dataset contexts (hospital readmission, loan default, customer churn, manufacturing defect, and employee promotion prediction). Each feature carries a ground-truth risk level and leakage type. The benchmark includes 13 high-risk, 1 medium-risk, and 16 safe features, covering temporal, post-hoc, and proxy leakage patterns. We evaluate at two operating thresholds—“high” (flag only high-risk features) and “medium” (flag medium and above)—and report binary precision, recall, and F_1 .

Table 2. Semantic leakage module evaluation on 30-feature benchmark.

Threshold	P	R	F_1	TP	FP	FN
high	1.000	1.000	1.000	13	0	0
medium	1.000	0.929	0.963	13	0	1

The single false negative at the medium threshold is `treatment_count`, a genuinely ambiguous feature: it may represent prior treatment history (a legitimate predictor) or treatments administered during the current hospital stay (a post-hoc quantity). Without explicit dataset documentation, linguistic reasoning alone is insufficient to disambiguate this case—a limitation acknowledged in Section 7. The benchmark and evaluation script are included in the repository at `data/semantic_benchmark.json` and `scripts/evaluate_semantic_leakage.py`.

5 Experimental Evaluation

5.1 Datasets

We evaluate on eleven datasets: six synthetic and five real-world. **Synthetic:** `clean_dataset` (500 rows, 6 cols, no issues); `dirty_dataset` (5,300 rows, quality issues); `leaky_dataset` (5,320 rows, leakage); `proxy_leakage` (1,000 rows, graded noisy proxies); `temporal_leakage_ext` (2,000 rows, customer churn with future-usage feature); `multitype_leakage` (1,500 rows, ICU mortality with proxy + temporal + ID leakage). **Real-world:** Titanic (1,309 rows), Pima Diabetes (768), Adult Census Income (48,842), German Credit (1,000), and Wine Quality Red (1,599)—all obtained via OpenML.

5.2 Readiness Score Results

Table 3 reports readiness scores for the five principal evaluation datasets. The clean control dataset achieves grade A (95/100), confirming zero false positives, while datasets with leakage and quality issues are appropriately downgraded to B/C. Figure 2 visualises these scores alongside the grade thresholds.

Table 3. Readiness scores on five datasets (21 checks each).

Dataset	Score	Grade	Pass	Key findings
<code>clean_dataset</code>	95.0	A	21/21	Zero false positives
<code>dirty_dataset</code>	65.0	C	12/21	5 quality, 2 sufficiency
<code>leaky_dataset</code>	72.0	B	11/21	Proxy, temporal, ID
Titanic	75.0	B	17/21	boat $r=0.97$
Pima Diabetes	80.0	B	20/21	51 outliers

5.3 Extended UCI Evaluation

Table 4 reports readiness scores on four additional real-world datasets obtained from OpenML, demonstrating the framework’s behaviour across different domains, sizes, and feature types. The Adult Census




Figure 2. Readiness scores across five datasets. The clean control achieves grade A (95/100) with zero false positives. Dashed lines mark grade thresholds; all real-world and dirty datasets land in the B–C range.

Income dataset (48,842 rows) scores 68/100 (C), driven primarily by significant class imbalance (income >50K accounts for only 23.9% of records) and nine mixed-type categorical features with high missing rates in `workclass` and `occupation`. German Credit scores 74/100 (B): the dataset is well-curated by construction but the 70/30 class split triggers a class imbalance warning, and several categorical features exhibit high within-class correlation. Heart Disease (303 rows) scores only 62/100 (C), with the sufficiency checks reporting $n/p = 21.6$ (below the comfortable threshold of 50) and high inter-feature correlation between `thalach` and `age`. Wine Quality achieves the highest score among the UCI group at 82/100 (B), with no leakage detected and well-behaved feature distributions; the main penalty arises from class imbalance (quality ≥ 6 accounts for 53.5% of samples, but the minority class in the binarised target is adequately supported).

Table 4. Readiness scores on four additional UCI/OpenML datasets.

Dataset	Score	Grade	Pass	Primary issues detected
Adult Census	68	C	13/21	Class imbalance (76/24), missing values in 2 cols, low MI relevance in <code>fnlwgt</code>
German Credit	74	B	16/21	Class imbalance (70/30), high correlation between credit features
Heart Disease	62	C	11/21	Low $n/p = 21.6$, inter-feature correlation, 6 outlier cols
Wine Quality	82	B	18/21	Mild class imbalance, skewness in <code>residual.sugar</code>

5.4 Impact Analysis: Leakage Quantification

Table 5 reports cross-validated accuracy before and after removing leaky features. A baseline accuracy of 1.000 drops to ≈ 0.952 after cleaning ($\Delta \approx -0.048$), confirming that the original high accuracy was entirely attributable to leaked information.

Table 5. Impact analysis on `leaky_dataset`: 5-fold CV accuracy.

Model	Baseline	Cleaned	Δ
Logistic Regression	1.000 ± 0.000	0.952 ± 0.008	-0.048
Random Forest	1.000 ± 0.000	0.954 ± 0.010	-0.046
XGBoost	1.000 ± 0.000	0.951 ± 0.009	-0.049

5.5 Unified Leakage Risk Score Validation

Table 6 validates Eq. (1) on representative features. All intentionally leaky features achieve $\mathcal{L} \geq 0.83$ while clean features remain below 0.30.

Table 6. Unified leakage risk scores $\mathcal{L}(f)$ per component.

Feature	Type	ρ	$\tilde{I}$	π	$\mathcal{L}$
proxy (corr=1.0)	leaky	1.00	1.00	1.00	1.000
noisy proxy ($r \approx 0.97$)	leaky	0.97	0.93	0.91	0.940
indirect feature	leaky	0.82	0.79	0.88	0.831
age (unrelated)	clean	0.04	0.07	0.06	0.057
income (moderate)	clean	0.28	0.31	0.24	0.281

5.6 Real-World Case Studies

Beyond aggregate scores, we examined three real-world datasets in detail to illustrate how the framework identifies specific problematic features in practice (Table 7).

Table 7. Real-world case study summary (21 checks, impact analysis skipped).

Dataset	Rows	Score	Grade	Primary findings
Titanic	1,309	80.6	B	Target leakage (boat, $r=0.97$); 4 cols with >5% missing; 99 outliers
Adult Census	48,842	92.4	A	2 cols with missing values; 6,854 outliers (skewed capital-gain); 52 duplicates
German Credit	1,000	97.5	A	3 near-zero MI features; 25 outliers

On Titanic, the framework flags boat ($r=0.97$, $\mathcal{L}=0.94$) as a post-hoc leaky feature (lifeboat assignment was recorded after the outcome) alongside four columns exceeding the 5% missing-value threshold and a distributional shift between dataset halves. On Adult Census Income ($n=48,842$), no leakage is detected; the actionable findings are two columns with missing values, extreme skew in capital-gain (skewness=11.9), and 52 duplicate rows. German Credit scores highest (97.5/A), with the only notable issue being three features with near-zero mutual information with the target—candidates for removal consistent with prior analysis of this dataset [?]. These case studies demonstrate that the framework behaves sensibly across datasets with qualitatively different defect profiles: it correctly distinguishes leakage (Titanic), mild quality issues (Adult Census), and clean-but-low-signal features (German Credit).

6 Quantitative Benchmark

6.1 Feature Coverage

Figure 3 compares ml-framework against three actively maintained tools on 29 checks. ml-framework is the only tool to provide unified leakage risk scoring, temporal ordering verification, statistical sufficiency checks, and actionable code-level recommendations.

6.2 Leakage Detection Rate

Table 8 reports results on four constructed scenarios. ml-framework detects all four; Great Expectations detects one (ID-column, via manual rule authoring); ydata-profiling and Deepchecks detect none.

Scope of comparison. It is important to note that Great Expectations, Deepchecks, and ydata-profiling were not designed as dedicated leakage detection frameworks. Great Expectations is primarily a data contract and expectations engine; ydata-profiling is an exploratory data analysis tool; Deepchecks focuses on model and dataset integrity for production ML systems. The comparison above is therefore best interpreted as an *ecosystem-level functionality audit* rather than a direct algorithmic competition: it documents which ML-relevant checks are available out-of-the-box in each tool, measured against the 29-check reference set that reflects what a thorough pre-training data review should cover. In their respective design spaces—data




Figure 3. Feature coverage comparison: number of checks supported out of 29. ml-framework provides $3.2\times$ more checks than the closest alternative (Deepchecks, 11/29).

Table 8. Leakage detection: ✓ = detected, ✗ = missed.

Scenario	ml-fw	ydata	DC	GE
Perfect proxy ($r=1.0$)	✓	✗	✗	✗
Noisy proxy ($r\approx 0.97$)	✓	✗	✗	✗
ID column (card. 100%)	✓	✗	✗	✓
Indirect computed feature	✓	✗	✗	✗
Detection rate	4/4	0/4	0/4	1/4

contracts, EDA, and model monitoring—the compared tools excel. ml-framework complements rather than replaces them, addressing the specific gap of automated, pre-training leakage detection with a calibrated readiness score.

6.3 Output and Configuration Flexibility

Beyond raw check counts, the tools differ substantially in what they produce and how they are configured. Table 9 summarises five qualitative dimensions. ml-framework is the only tool that generates an interpretable numerical readiness score, provides code-level corrective suggestions, and supports both a YAML-driven configuration and a plugin API for custom checks. Great Expectations offers the most flexible data contract system but requires manual expectation authoring and has no ML-specific output. ydata-profiling and Deepchecks produce rich HTML reports but offer no per-dataset score or automated recommendations.

Table 9. Output and configuration flexibility comparison.

Capability	ml-fw	ydata	DC	GE
Numerical readiness score	✓	✗	✗	✗
Letter grade (A–F)	✓	✗	✗	✗
Code-level recommendations	✓	✗	✗	✗
YAML / file-based config	✓	✗	✗	✓
Custom check plugin API	✓	✗	✗	✓
HTML report export	✓	✓	✓	✓
JSON export	✓	✗	✓	✓
Interactive web UI	✓	✗	✗	✓
Python SDK	✓	✗	✓	✓
CLI	✓	✗	✗	✓

6.4 Runtime Comparison

Table 10 reports wall-clock time on three dataset sizes. ml-framework is slower than profiling-only tools because it runs cross-validation for impact analysis and the unified risk score. This overhead is justified by

the much richer analytical output. Deepchecks produced runtime errors on Python 3.13 due to a `np.Inf` API removal in NumPy 2.0.

Table 10. Runtime (seconds). “—” = compatibility error with NumPy 2.0.

Dataset	ml-fw	ydata	DC	GE
500 rows	2.88	0.15	—	0.01
1,000 rows	2.23	0.15	—	0.01
5,300 rows	7.54	0.16	—	0.01

7 Discussion

Strengths. ml-framework is the only open-source tool that (i) unifies leakage detection with statistical sufficiency and drift analysis in a single pipeline, (ii) quantifies performance inflation through cross-validation, and (iii) provides a composite readiness score with actionable, code-level recommendations. The unified leakage risk score (Eq. 1) addresses the limitation that individual heuristics are insufficiently principled from a research perspective, by combining three complementary signals grounded in information theory.

A key empirical finding is that the three components of $\mathcal{L}(f)$ are genuinely complementary. For the Titanic boat column, the correlation component alone ($\rho = 0.97$) would be sufficient to flag the feature. However, for indirect features where the relationship is nonlinear (e.g. a binary target encoded via a monotone but nonlinear transformation), Pearson $|r|$ can be as low as 0.50 while $\tilde{I}$ and π remain above 0.80. Conversely, for high-cardinality categorical features, Cramér’s V is a more reliable signal than MI. The weighted combination reduces false negatives compared to any single check applied in isolation.

The semantic leakage module adds a qualitatively different capability: it can flag features like `discharge_code` (a post-hoc administrative code assigned at patient discharge) or `future_usage` (measured after the prediction window) with zero false positives in our experiments, relying on linguistic reasoning that is simply outside the scope of statistical tests. The GPT-4o-mini model correctly identified all high-risk features in the `multitype_leakage` synthetic dataset and produced explanations that are directly usable in a data documentation workflow.

Limitations. The framework currently handles only binary and multi-class classification; regression tasks require a modified performance-inflation metric (e.g. coefficient of determination R^2 delta rather than accuracy delta). Time-series tasks with auto-regressive features require special treatment, as the temporal split logic must account for forecast horizons rather than simply checking monotonic date ordering.

The per-feature CV in the unified risk score adds $O(p)$ computational overhead, which becomes noticeable on wide datasets ($p > 500$). In practice, this can be mitigated by limiting the risk score to features flagged by the cheaper correlation check as a first pass, or by reducing the number of CV folds.

The LLM semantic module introduces a cloud API dependency, latency (typically 5–15 seconds for 30 features), and cost (approximately \$0.001 per dataset call at GPT-4o-mini pricing). The quality of semantic analysis is also sensitive to the user-provided dataset description: brief or uninformative descriptions reduce the model’s ability to distinguish temporal leakage from legitimate predictive features. Finally, the module cannot guarantee completeness—unusual naming conventions or domain-specific abbreviations may evade detection.

Future Work. We identify four directions for extending ml-framework. First, support for regression and time-series tasks requires new performance-inflation metrics and temporally-aware sufficiency checks. Second, causal discovery methods—such as the PC algorithm or FCI—could detect indirect leakage paths that information-theoretic measures miss when the causal graph is non-trivial. Third, while we provide an initial 30-feature semantic leakage benchmark (Section 4), a larger benchmark derived from real data

quality incident reports would enable more statistically robust evaluation of the LLM module across diverse domain vocabularies and naming conventions. Fourth, the readiness score weighting (w_q, w_l, w_f, w_s) is currently fixed; a data-driven calibration approach that learns weights from labelled examples of production failures would make the score more adaptive to domain-specific risk profiles.

8 Conclusion

We presented **ml-framework**, a comprehensive open-source toolkit for automated dataset quality assessment and data leakage detection. The 21-check pipeline, novel unified leakage risk score, and LLM-assisted semantic analysis advance the state of the art beyond heuristic thresholds. Empirical evaluation on eleven datasets shows 4/4 leakage detection rate, calibrated readiness scores from 65 (C, dirty data) to 95 (A, clean data), and coverage of $3.2\times$ more checks than the best competing tool. We release all code, datasets, benchmarks, and the Streamlit demo at <https://github.com/jaimecano12/ml-framework>.

Acknowledgements

The author thanks his supervisor for guidance on experimental design and the suggestion to incorporate LLM-based semantic leakage analysis.

References

- [1] S. Kaufman, S. Rosset, C. Perlich, and O. Stitelman, “Leakage in data mining: Formalization, detection, and avoidance,” *ACM Transactions on Knowledge Discovery from Data*, vol. 6, no. 4, pp. 1–21, 2012.
- [2] D. Sculley et al., “Hidden technical debt in machine learning systems,” in *Advances in NeurIPS*, vol. 28, 2015.
- [3] E. Breck, M. Zinkevich, N. Polyzotis, S. Whang, and S. Roy, “Data validation for machine learning,” in *Proc. 2nd SysML Conference*, 2019.
- [4] S. D. F. M. van den Burg and T. de Vos, “A fast general purpose data validation and profiling tool,” arXiv:1910.15327, 2019.
- [5] A. Campbell and B. Ulery, “Great Expectations: Always know what to expect from your data,” <https://greatexpectations.io>, 2019.
- [6] Deepchecks Team, “Deepchecks: Testing and validating ML models and data,” <https://deepchecks.com>, 2021.
- [7] F. Pedregosa et al., “Scikit-learn: Machine learning in Python,” *JMLR*, vol. 12, pp. 2825–2830, 2011.
- [8] T. Chen and C. Guestrin, “XGBoost: A scalable tree boosting system,” in *Proc. ACM KDD*, pp. 785–794, 2016.
- [9] A. Kraskov, H. Stögbauer, and P. Grassberger, “Estimating mutual information,” *Physical Review E*, vol. 69, no. 6, 2004.
- [10] B. C. Ross, “Mutual information between discrete and continuous data sets,” *PLOS ONE*, vol. 9, no. 2, 2014.
- [11] S. Rabanser, S. Günnemann, and I. Steinwart, “Failing loudly: An empirical study of methods for detecting dataset shift,” in *Advances in NeurIPS*, vol. 32, 2019.
- [12] N. Siddiqi, *Credit Risk Scorecards*. John Wiley & Sons, 2006.

- [13] W. McKinney, “Data structures for statistical computing in Python,” in *Proc. 9th Python in Science Conference*, pp. 56–61, 2010.
- [14] OpenAI, “GPT-4 technical report,” arXiv:2303.08774, 2023.
- [15] D. Dua and C. Graff, “UCI Machine Learning Repository,” UC Irvine, 2017. <https://archive.ics.uci.edu/ml>
- [16] K. Zha et al., “Data-centric artificial intelligence: A survey,” arXiv:2303.10158, 2023.
- [17] S. Narayan, M. Hearst, A. Halevy, and P. Haas, “Can foundation models wrangle your data?” arXiv:2205.09911, 2022.
- [18] Y. Sui, J. Zhou, M. Zhou, S. Han, and J. Zhang, “TAP4LLM: Table provider on sampling, augmenting, and packing semi-structured data for large language model reasoning,” arXiv:2312.09039, 2023.
- [19] A. Ng, “A chat with Andrew on MLOps: From model-centric to data-centric AI,” DeepLearning.AI, 2021. <https://www.youtube.com/watch?v=06-AZXmwHjo>